

MLRF Lecture 02

J. Chazalon, LRDE/EPITA, 2019

Global image descriptors

Lecture 02 part 02

Two approaches

Lecture 4

Global image descriptors

- Compute **statistics about the content** of the image
- Produce a **single global vector**

Very attractive because they are very fast to compute and match, but... (see end of section)

Bag of Features techniques

- **Select regions** of interest in the image (may be a variable quantity)
- **Compute descriptors** for each region
- **Index each part** separately (like a text search engine which indexes words)

It is always possible to build a single descriptor from local descriptors!

This technique is the one used in modern image search engines.

Color Histograms

Color histograms – a very simple global descriptor (of pixels statistics)

High invariance to many transformation

rotation, scaling thanks to normalization, perspective...


But limited discriminative power

Easy to implement


1. Reduce the colors (opt. when performing backprojection)
2. Compute a reduced color histogram on each image
3. Use a distribution distance to compare the descriptors

Color histograms: Some results on *Twin it!*


22 86:0.005 257:0.007 156:0.008 13:0.009




23 378:0.012 320:0.019 297:0.037 263:0.040




24 331:0.024 302:0.035 323:0.035 271:0.042




25 376:0.038 197:0.043 66:0.046 205:0.056




34 318:0.020 42:0.028 242:0.041 102:0.042




35 219:0.002 335:0.005 362:0.024 251:0.034



36 69:0.025 155:0.028 304:0.030 212:0.034



37 308:0.003 108:0.011 365:0.031 351:0.034



Timing comparison (1 CPU)

Template matching

Match each pair of image:
3 hours

Color Histogram

Color reduction:
3 seconds

Compute color histogram for
all bubbles:
30 seconds

Compute distance between
each pair of descriptors:
2 seconds

Color histograms: Step by step

1: Color reduction

1. Use K-Means or any other clustering technique to find N useful colors.
2. Project each pixel value on the value of the closest cluster center.

Swain & Ballard 1991




Fig. 1. Left: Image of a Crunchberries cereal box. Right: Three dimensional color histogram of the Crunchberries image with the black background substrated.



↑ 16M colors





↑ 7 colors (+ white bg)

One possible result on the *Twin it!* poster

Color histograms: Step by step

2: Histogram computation

You already know it.
(*Normalize it.*)

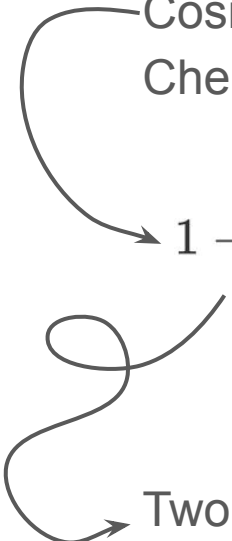


Color histograms: Step by step

3: Descriptor comparison

Many distribution metrics.

Cosine, Euclidean,
Chebyshev...


$$1 - \frac{u \cdot v}{||u||_2 ||v||_2}.$$

Read, try, compare, learn!

Two histograms =
Two 1-D vectors!

```
from scipy.spatial.distance import ...
```

Distance functions between two numeric vectors `u` and `v`. Computing distances over a large collection of vectors is inefficient for these functions. Use `pdist` for this purpose.

<code>braycurtis(u, v[, w])</code>	Compute the Bray-Curtis distance between two 1-D arrays.
<code>canberra(u, v[, w])</code>	Compute the Canberra distance between two 1-D arrays.
<code>chebyshev(u, v[, w])</code>	Compute the Chebyshev distance.
<code>cityblock(u, v[, w])</code>	Compute the City Block (Manhattan) distance.
<code>correlation(u, v[, w, centered])</code>	Compute the correlation distance between two 1-D arrays.
<code>cosine(u, v[, w])</code>	Compute the Cosine distance between 1-D arrays.
<code>euclidean(u, v[, w])</code>	Computes the Euclidean distance between two 1-D arrays.
<code>jensenshannon(p, q[, base])</code>	Compute the Jensen-Shannon distance (metric) between two 1-D probability arrays.
<code>mahalanobis(u, v, VI)</code>	Compute the Mahalanobis distance between two 1-D arrays.
<code>minkowski(u, v[, p, w])</code>	Compute the Minkowski distance between two 1-D arrays.
<code>seuclidean(u, v[, w])</code>	Return the standardized Euclidean distance between two 1-D arrays.
<code>squeuclidean(u, v[, w])</code>	Compute the squared Euclidean distance between two 1-D arrays.
<code>wminkowski(u, v[, p, w])</code>	Compute the weighted Minkowski distance between two 1-D arrays.

Discussion

Can you think of other global descriptors we could have implemented for the *Twin it!* case?

Other global image descriptors

More global descriptors

GIST of a scene:

- Oliva, Torralba, “Modeling the shape of the scene: a holistic representation of the spatial envelope”, IJCV’01.
- Douze, Jegou, Sandhawalia, Amsaleg, Schmid, “Evaluation of GIST descriptors for web-scale image search”, CIVR’09.

CENTRIST: CENsus Transform hISTogram

- Wu, Rehg, “CENTRIST: a visual descriptor for scene categorization”, TPAMI’11.




Figure 15. Organization of natural scenes according to the openness and ruggedness properties estimated by the WDSTs.

Global descriptors: drawback

According to F. Perronnin:

Highly efficient to compute and to match

⇒ **perfect in theory**

But **robustness vs informativeness tradeoff is hard to set**

(personal conclusion):

- Approaches based on **global image descriptors** are confined to **near-duplicate detection** applications until now.
- Modern search engines use local representations and leverage them.